

STATS 3DA3

Homework Assignment 6

JC Abanto (400377157) , Zhiyan Chen (400365265), Jasmine Ho (400374898)

2025-04-16

Question

1)

Our target variable measure the severity of heart disease which can be defined as binary classification problem. Our goal would then to be able to predict the probability of heart disease. This means we can carry out a logistic regression and a random forest classifier to predict the presence and absence of heart disease.

2)

```
from ucimlrepo import fetch_ucirepo
from sklearn.preprocessing import StandardScaler

heart_disease = fetch_ucirepo(id=45)

X = heart_disease.data.features
y = heart_disease.data.targets

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X)
```

3)

We have 13 features and 1 target variable. Starting with our features, we have age (age in years), sex (1 = male; 0 = female), cp (1: typical angina, 2: atypical angina, 3: non-anginal pain, 4: asymptomatic), trestbps (resting blood pressure in mm Hg on admission to the hospital), chol (serum cholestoral in mg/dl), fbs (fasting blood sugar > 120 mg/dl where 1 = true; 0 = false), restecg (resting electrocardiographic results), exang (exercise induced angina), oldpeak (ST depression induced by exercise relative to rest), slope (1 = upsloping, 2 = flat, 3 = downsloping), ca (number of major vessels colored by floursopy), thal (3 = normal; 6 = fixed defect; 7 = reversable defect). Finally our target variable, 'num', is the diagnosis of heart disease.

```
print(f"Observations in X: {len(X)}")
print(f"Summary of X:\n{X.describe()}")
print(f"Summary of y:\n{y.describe()}")
```

Observations in X: 303

Summary of X:

	age	sex	cp	trestbps	chol	fbs \
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	54.438944	0.679868	3.158416	131.689769	246.693069	0.148515
std	9.038662	0.467299	0.960126	17.599748	51.776918	0.356198
min	29.000000	0.000000	1.000000	94.000000	126.000000	0.000000
25%	48.000000	0.000000	3.000000	120.000000	211.000000	0.000000
50%	56.000000	1.000000	3.000000	130.000000	241.000000	0.000000
75%	61.000000	1.000000	4.000000	140.000000	275.000000	0.000000
max	77.000000	1.000000	4.000000	200.000000	564.000000	1.000000

	restecg	thalach	exang	oldpeak	slope	ca \
count	303.000000	303.000000	303.000000	303.000000	303.000000	299.000000
mean	0.990099	149.607261	0.326733	1.039604	1.600660	0.672241
std	0.994971	22.875003	0.469794	1.161075	0.616226	0.937438
min	0.000000	71.000000	0.000000	0.000000	1.000000	0.000000
25%	0.000000	133.500000	0.000000	0.000000	1.000000	0.000000
50%	1.000000	153.000000	0.000000	0.800000	2.000000	0.000000
75%	2.000000	166.000000	1.000000	1.600000	2.000000	1.000000
max	2.000000	202.000000	1.000000	6.200000	3.000000	3.000000

	thal
count	301.000000
mean	4.734219
std	1.939706
min	3.000000
25%	3.000000

50%	3.000000
75%	7.000000
max	7.000000

Summary of y:

	num
count	303.000000
mean	0.937294
std	1.228536
min	0.000000
25%	0.000000
50%	0.000000
75%	2.000000
max	4.000000

We find that the average age of patients is 54.4 years old, with a standard deviation of 9.1 years.

```
print(f"Data Types of X: \n{X.dtypes}")
print(f"\nData Types of y: \n{y.dtypes}")
```

Data Types of X:

age	int64
sex	int64
cp	int64
trestbps	int64
chol	int64
fbs	int64
restecg	int64
thalach	int64
exang	int64
oldpeak	float64
slope	int64
ca	float64

```
thal          float64
dtype: object
```

Data Types of y:

```
num          int64
dtype: object
```

All of our data types in X are numerical but some representing categorical variables.

4)

```
print(f"y before transformation: {y['num'].value_counts()}")
y['num'] = y['num'].apply(lambda x: 1 if x > 0 else 0)
print(f"y before transformation: {y['num'].value_counts()}")
```

```
y before transformation: num
```

```
0    164
1     55
2     36
3     35
4     13
```

```
Name: count, dtype: int64
```

```
y before transformation: num
```

```
0    164
1    139
```

```
Name: count, dtype: int64
```

```
/var/folders/5q/w2pw9gp90mldwv2xvm9ml_jc0000gn/T/ipykernel_9942/2854447611.py:2: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/

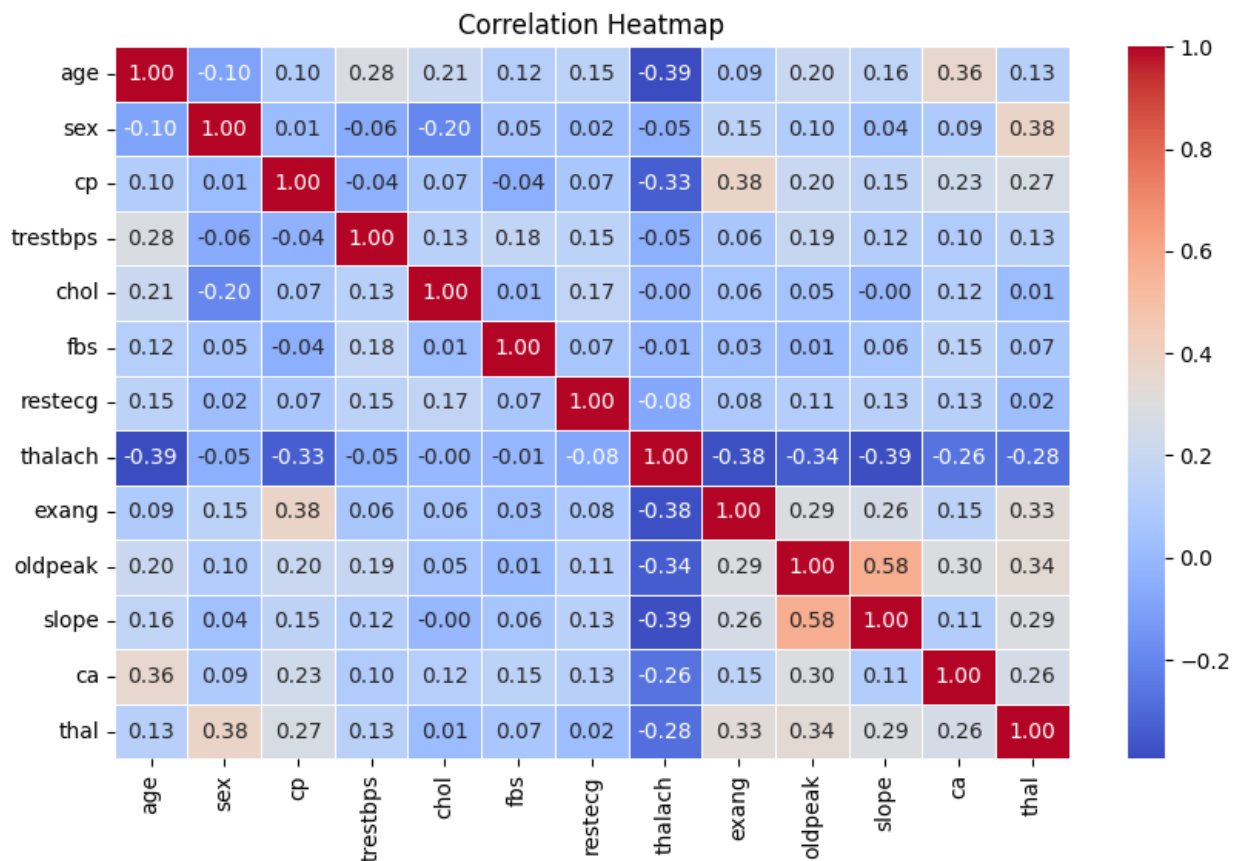
```
y['num'] = y['num'].apply(lambda x: 1 if x > 0 else 0)
```

5)

```
import seaborn as sns
import matplotlib.pyplot as plt

corr_matrix = X.corr()

plt.figure(figsize=(10, 6))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=0.5)
plt.title("Correlation Heatmap")
plt.show()
```



From this correlation matrix, we can conclude that thalach (max heart rate) has a strong negative correlation (-0.39) with age and oldpeak. We can assume that younger individuals tend to have a higher heart rate, while those with more severe heart disease (higher oldpeak) have lower thalach.

We also found that ca (number of major vessels) has a strong positive correlation (0.36) with age. We can say that older individuals are more likely to have more blocked vessels.

6)

```
X = X.dropna()
y = y.loc[X.index]

print(f"Length of X after transformations: {len(X)}")
print(f"Length of y after transformations: {len(y)}")
```

Length of X after transformations: 297

Length of y after transformations: 297

7)

```
import pandas as pd
from sklearn.metrics import silhouette_samples, silhouette_score
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

categorical_columns = ['sex', 'cp', 'fbs', 'restecg', 'exang', 'slope', 'thal']
X_cleaned = X.drop(columns=categorical_columns)

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_cleaned)
```

```

k_range = range(2,10)
silhouette_scores = []

for k in k_range:
    kmeans = KMeans(n_clusters = k, n_init = 20, random_state = 0)
    cluster_labels = kmeans.fit_predict(X_scaled)
    silhouette_avg = silhouette_score(X_scaled, cluster_labels)
    silhouette_scores.append(silhouette_avg)

plt.figure(figsize=(9, 5))
plt.plot(k_range, silhouette_scores, marker='o')
plt.ylabel("Silhouette Score")
plt.xlabel("Number of Clusters (k)")
plt.title("k Values Silhouette Score")
plt.show()

```




```

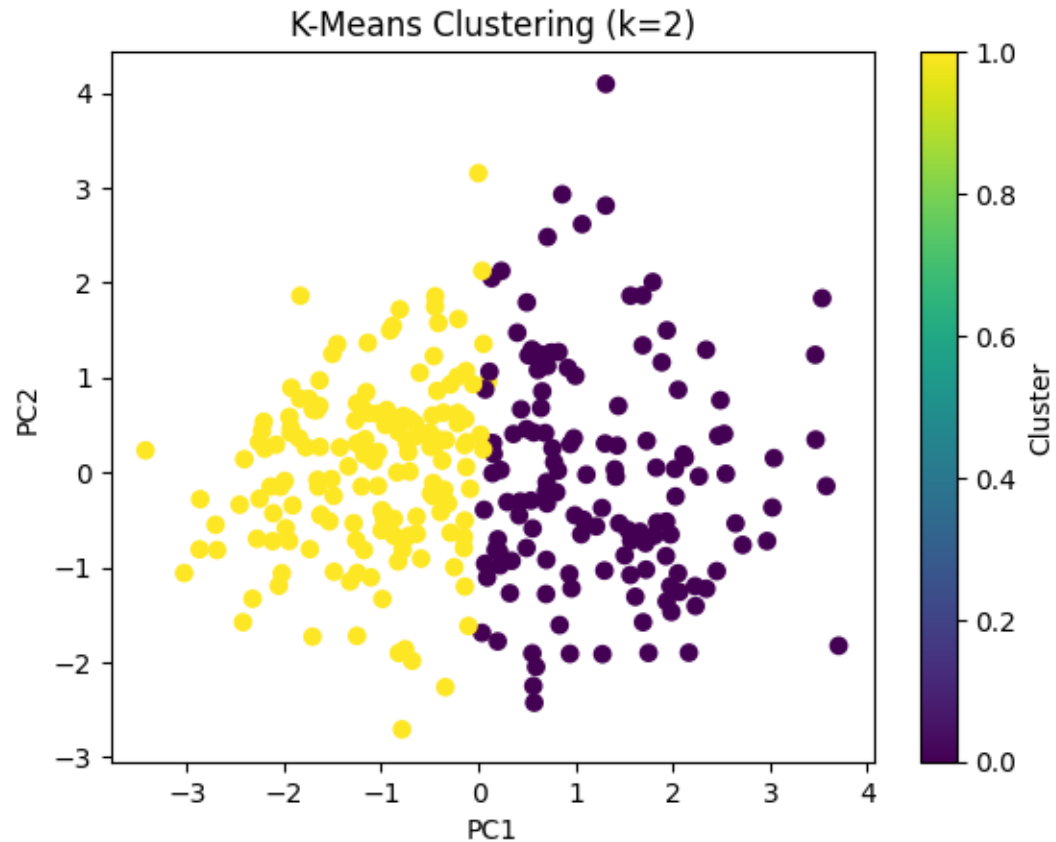
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)

kmeans = KMeans(n_clusters=2, random_state=1)
clusters = kmeans.fit_predict(X_scaled)

plt.scatter(X_pca[:, 0], X_pca[:, 1], c=clusters, cmap='viridis')
plt.title("K-Means Clustering (k=2)")
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.colorbar(label="Cluster")
plt.show()

X_scaled_df = pd.DataFrame(X_scaled, columns=X_cleaned.columns)
X_scaled_df['cluster'] = clusters

```



8)

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, random_state=1)
```

9)

We are going to use logistic regression and random forest. Logistic regression is suitable for this assignment because it predicts a binary outcome, in which case, the target variable is 0 or 1. Random forest is also a good classifier to use because it has high predictive accuracy and does not depend on linear relationships, which is good for this dataset as there are both numerical and categorical variables.

10)

We are going to use accuracy and F1 scores to compare the classifier performance between logistic regression and random forest. We can create the confusion matrix from the predictions to calculate the accuracy and F1 scores. Accuracy scores are calculated by the number of correct predictions over the total number of predictions.

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

Where:

TP = True Positives

TN = True Negatives

FP = False Positives

FN = False Negatives

To obtain the F1 score, we will need to use precision and recall that are derived from the confusion matrix. Once we calculate that, the F1 score can be calculated by:

$$\text{Precision} = \frac{TP}{TP + FP}$$

$$\text{Recall} = \frac{TP}{TP + FN}$$

$$F_1 = \frac{2 \times \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

Using these metrics, we can compare the overall accuracy and balance between the two classifiers to determine which classifier most optimal for predicting heart disease.

11)

Logistic Regression

```
import numpy as np
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

import warnings
warnings.filterwarnings('ignore')

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

log_reg = LogisticRegression(solver='liblinear', max_iter=1000, random_state=42)

param_grid = {
    'penalty': ['l1', 'l2'],
```

```

    'C': [0.001, 0.01, 0.1, 1, 10, 100, 1000] # test this out to see if accuracy improves
}

grid_search = GridSearchCV(estimator=log_reg, param_grid=param_grid, cv=5, scoring='accuracy',
grid_search.fit(X_train_scaled, y_train)

print("Optimal parameters: ", grid_search.best_params_)
print("Cross-validation accuracy: {:.4f}".format(grid_search.best_score_))

best_log_reg = grid_search.best_estimator_
y_pred = best_log_reg.predict(X_test_scaled)
y_pred2 = best_log_reg.predict(X_train_scaled)

accuracy = accuracy_score(y_test, y_pred)
train_accuracy = accuracy_score(y_train, y_pred2)
cm = confusion_matrix(y_test, y_pred)
report = classification_report(y_test, y_pred, target_names=['No Heart Disease', 'Heart Disease'])

print(f"Test set accuracy: {accuracy:.4f}")
print(f"Train set accuracy: {train_accuracy:.4f}")
print("Confusion Matrix:")
print(cm)
print("Classification Report:")
print(report)

```

Fitting 5 folds for each of 14 candidates, totalling 70 fits

Optimal parameters: {'C': 0.01, 'penalty': 'l2'}

Cross-validation accuracy: 0.8456

Test set accuracy: 0.8556

Train set accuracy: 0.8454

Confusion Matrix:

```
[[45  6]
```

[7 32]]

Classification Report:

	precision	recall	f1-score	support
No Heart Disease	0.87	0.88	0.87	51
Heart Disease	0.84	0.82	0.83	39
accuracy			0.86	90
macro avg	0.85	0.85	0.85	90
weighted avg	0.86	0.86	0.86	90

Random Forest

```
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

import warnings
warnings.filterwarnings('ignore')

rf_clf = RandomForestClassifier(random_state=42)
rf_param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'max_features': ['sqrt', 'log2', None],
    'criterion': ['gini', 'entropy']
}

rf_grid_search = GridSearchCV(estimator=rf_clf, param_grid=rf_param_grid, cv=5, scoring='accuracy')
```

```

rf_grid_search.fit(X_train_scaled, y_train)

print("Optimal Parameters:", rf_grid_search.best_params_)
print("Cross-Validation Accuracy:", rf_grid_search.best_score_)

best_rf_clf = rf_grid_search.best_estimator_
rf_y_pred = best_rf_clf.predict(X_test_scaled)

rf_accuracy = accuracy_score(y_test, rf_y_pred)
rf_cm = confusion_matrix(y_test, rf_y_pred)
rf_report = classification_report(y_test, rf_y_pred, target_names=['No Heart Disease', 'Heart Disease'])

print(f"Test Accuracy: {rf_accuracy:.4f}")
print("Confusion Matrix:")
print(rf_cm)
print("Classification Report:")
print(rf_report)

```

Fitting 5 folds for each of 216 candidates, totalling 1080 fits

Optimal Parameters: {'criterion': 'gini', 'max_depth': None, 'max_features': 'sqrt', 'min_samples_split': 2, 'min_samples_leaf': 1}

Cross-Validation Accuracy: 0.8166085946573751

Test Accuracy: 0.8444

Confusion Matrix:

```

[[44  7]
 [ 7 32]]

```

Classification Report:

	precision	recall	f1-score	support
No Heart Disease	0.86	0.86	0.86	51
Heart Disease	0.82	0.82	0.82	39
accuracy			0.84	90

macro avg	0.84	0.84	0.84	90
weighted avg	0.84	0.84	0.84	90

12)

We are going to apply the LASSO for feature selection to our logistic regression classifier.

```
print(X.shape)
print(y.shape)
print(X_train_scaled.shape)
```

```
(297, 13)
```

```
(297,)
```

```
(207, 14)
```

Note: When training the LASSO, we realized that X (full feature set) has 13 columns, but LogisticRegression (scaled training data) is expecting 14 features as input. Therefore, we want to make sure the target column is not in X.

```
# drop target column
df = pd.concat([X, y], axis=1)
X = df.drop(columns=['num'])
y = df['num']
```

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=1)

from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test) # scale only X
```

```
print(X_train_scaled.shape)
```

(207, 13)

Now we can finally train the LASSO.

```
from sklearn.linear_model import LogisticRegressionCV, LogisticRegression
from sklearn.metrics import accuracy_score, roc_curve, roc_auc_score
import pandas as pd
import numpy as np

#lasso with cross-validation
lasso_cv = LogisticRegressionCV(
    penalty='l1',
    solver='saga',
    cv=5,
    max_iter=10000,
    scoring='accuracy',
    random_state=42
)
lasso_cv.fit(X_train_scaled, y_train)
```

```
LogisticRegressionCV(cv=5, max_iter=10000, penalty='l1', random_state=42,
                      scoring='accuracy', solver='saga')
```

```
#find the best c and lambda
best_c = lasso_cv.C_[0]
print("Best C (1/lambda):", best_c)
print("Best lambda (1/C):", 1 / best_c)
```

Best C (1/lambda): 2.782559402207126

Best lambda (1/C): 0.3593813663804626


```
# train lasso again with the best C
m_lasso = LogisticRegression(
    penalty='l1',
    C=best_c,
    solver='saga',
    max_iter=10000,
    random_state=42
)
m_lasso.fit(X_train_scaled, y_train)
```

```
LogisticRegression(C=2.782559402207126, max_iter=10000, penalty='l1',
                    random_state=42, solver='saga')
```

```
# selected features with coefficients
coeffs = m_lasso.coef_.reshape(-1)
feature_names = X.columns
pd.DataFrame({
    'Feature': feature_names,
    'Coefficient': coeffs
})
```

	Feature	Coefficient
0	age	-0.054183
1	sex	0.893613
2	cp	0.560257
3	trestbps	0.636381
4	chol	0.309335
5	fbs	-0.666898
6	restecg	0.314907
7	thalach	-0.458236
8	exang	0.464750

	Feature	Coefficient
9	oldpeak	0.491677
10	slope	0.289350
11	ca	1.342414
12	thal	0.571055

```
# class 1 probabilities
```

```
m_lasso_pre_prob = m_lasso.predict_proba(X_test_scaled)[: , 1]
```

```
df_eval = pd.DataFrame({'prob': m_lasso_pre_prob, 'y_test': y_test})
```

```
# ROC and KS threshold
```

```
fpr, tpr, thresholds = roc_curve(df_eval.y_test, df_eval.prob)
```

```
ks_statistic = np.max(tpr - fpr)
```

```
ks_threshold = thresholds[np.argmax(tpr - fpr)]
```

```
print("KS threshold:", ks_threshold)
```

```
KS threshold: 0.42760641677437833
```

```
# sensitivity & specificity
```

```
ind = np.where(np.isclose(thresholds, ks_threshold, atol=0.001))
```

```
print("Sensitivity:", tpr[ind])
```

```
print("Specificity:", 1 - fpr[ind])
```

```
Sensitivity: [0.87179487]
```

```
Specificity: [0.76470588]
```

```
# classify based on KS threshold
```

```
df_eval['y_test_pred'] = df_eval.prob.map(lambda x: 1 if x > ks_threshold else 0)
```

```
# Accuracy, AUC and Confusion Matrix
print("KS Threshold Accuracy:", accuracy_score(df_eval.y_test, df_eval.y_test_pred))
print("AUC Score:", roc_auc_score(df_eval.y_test, df_eval.prob))
print("Confusion Matrix:", confusion_matrix(df_eval.y_test, df_eval.y_test_pred))
```

```
KS Threshold Accuracy: 0.8
AUC Score: 0.8798391151332328
Confusion Matrix: [[39 12]
 [ 6 33]]
```

```
# Classification Report for LASSO
lasso_report=classification_report(df_eval.y_test, df_eval.y_test_pred, target_names=['No Heart Disease', 'Heart Disease'])
print("Classification Report:")
print(lasso_report)
```

Classification Report:

	precision	recall	f1-score	support
No Heart Disease	0.87	0.76	0.81	51
Heart Disease	0.73	0.85	0.79	39
accuracy			0.80	90
macro avg	0.80	0.81	0.80	90
weighted avg	0.81	0.80	0.80	90

13)

Recall the three classifiers and their classification reports:

```
print("Logistic Regression Classification Report:")
print(report)
print("Random Forest Classification Report:")
```

```
print(rf_report)
print("LASSO Classification Report:")
print(lasso_report)
```

Logistic Regression Classification Report:

	precision	recall	f1-score	support
No Heart Disease	0.87	0.88	0.87	51
Heart Disease	0.84	0.82	0.83	39
accuracy			0.86	90
macro avg	0.85	0.85	0.85	90
weighted avg	0.86	0.86	0.86	90

Random Forest Classification Report:

	precision	recall	f1-score	support
No Heart Disease	0.86	0.86	0.86	51
Heart Disease	0.82	0.82	0.82	39
accuracy			0.84	90
macro avg	0.84	0.84	0.84	90
weighted avg	0.84	0.84	0.84	90

LASSO Classification Report:

	precision	recall	f1-score	support
No Heart Disease	0.87	0.76	0.81	51
Heart Disease	0.73	0.85	0.79	39
accuracy			0.80	90
macro avg	0.80	0.81	0.80	90

weighted avg	0.81	0.80	0.80	90
--------------	------	------	------	----

We will only focus on the two metrics we selected: **accuracy and F1 score**. By comparing them among the 3 classifiers, we have the following findings:

1. The **full logistic regression model** from (11) achieved the highest overall accuracy (0.86) and the highest heart disease F1 score (0.83), suggesting that it has excellent general performances.
2. While the **LASSO** has the lowest accuracy, it is still very acceptable (0.80). It also holds the highest recall (0.85) for predicting heart disease, meaning that it is the best model at identifying most of the true positive cases. (*recall's definition can be found in (10)*)

Impact of Feature Selection:

With the **LASSO**, we successfully reduced the numbers of features used in the model, keeping it interpretable. While the accuracy number went down a bit because of this simplification in the model, it maintained a very strong recall, which helped improve the model's performance on capturing positive heart disease cases. This is very useful and important for medical diagnostics dataset like this.

14)

From (13), we mentioned that while full logistic regression model had slightly higher accuracy and F1 score, the LASSO provided a more simplified and interpretable model by selecting only the most important features. For a medical diagnostics dataset like this, we prioritize the ability to catch more true heart disease cases over the highest possible accuracy - which is why we value a higher recall. We also prefer the simpler model as the reduction in complexity helps prevent overfitting, improves computational efficiency that are super essential for decision making. That is, we consider the **LASSO** as the best interpretable model.

```
# from (12), we had:
coeffs = m_lasso.coef_.reshape(-1)
feature_names = X.columns
features_coef = pd.DataFrame({
```

```

    'Feature': feature_names,
    'Coefficient': coeffs
})

# we want to make it in descending order to address the largest coef (most important feature)
important_features = features_coef[features_coef['Coefficient'] != 0].sort_values(by='Coefficient', ascending=False)

print("Important features by the LASSO:")
print(important_features)

```

Important features by the LASSO:

	Feature	Coefficient
11	ca	1.342414
1	sex	0.893613
3	trestbps	0.636381
12	thal	0.571055
2	cp	0.560257
9	oldpeak	0.491677
8	exang	0.464750
6	restecg	0.314907
4	chol	0.309335
10	slope	0.289350
0	age	-0.054183
7	thalach	-0.458236
5	fbs	-0.666898

From the output, we have the following findings:

1. The most important predictor variable found by the LASSO was **ca** (number of major vessels colored by fluoroscopy) since it has the largest positive coefficient. Note that we have standardized the train set, so this is in standardized units rather than the original scale of 0–3. This indicates that the more major vessels showing disease or blockage, the more likely an individual is to have heart disease.

2. There is a positive coefficient to **sex** (1 = male, 0 = female). This tells us that in this dataset, biological males are at a higher risk of heart disease. Note that some studies suggest that males do have a higher risk of severe heart attacks compared to women, but some also say this is a myth (Suman et al., 2023).
3. Another feature with high positive coefficient is **trestbps**, indicating that higher resting blood pressure is associated with a higher risk of heart disease.
4. **thalach** has a strong negative coefficient, meaning that individuals who achieves lower peak heart rates are more likely to have heart disease. This aligns with the clinical research result (Saxena et al., 2013).
5. **age** has a very small negative coefficient, which is surprising since there exist a lot of studies on the how heart disease risk increases with age, regardless of gender. However, this result only reflects that age is not the most dominant factor in this particular dataset.

15) [Bonus]

We are using the sub-groups from KMeans as a feature to improve the logistic regression model. After adding this feature and retraining the model, the classifier showed improved performance compared to the original model. We see that the test set accuracy improved from 0.8556 to 0.9000 and the F1-score also increased for both classes.

```
import warnings
warnings.filterwarnings('ignore')

categorical_columns = ['sex', 'cp', 'fbs', 'restecg', 'exang', 'slope', 'thal']
numerical_columns = [col for col in X.columns if col not in categorical_columns]
X_num = X[numerical_columns]

pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_num)

kmeans = KMeans(n_clusters=2, random_state=42)
```

```

clusters = kmeans.fit_predict(X_pca)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

X_train_num = X_train[numerical_columns]
X_test_num = X_test[numerical_columns]

X_train_pca = pca.transform(X_train_num)
X_test_pca = pca.transform(X_test_num)

train_clusters = kmeans.predict(X_train_pca).reshape(-1, 1)
test_clusters = kmeans.predict(X_test_pca).reshape(-1, 1)

X_train_num = np.hstack((X_train.values, train_clusters))
X_test_num = np.hstack((X_test.values, test_clusters))

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train_num)
X_test_scaled = scaler.transform(X_test_num)

log_reg = LogisticRegression(solver='liblinear', max_iter=1000, random_state=42)

param_grid = {
    'penalty': ['l1', 'l2'],
    'C': [0.001, 0.01, 0.1, 1, 10, 100, 1000]
}

grid_search = GridSearchCV(
    estimator=log_reg,
    param_grid=param_grid,
    cv=5,

```



```

        scoring='accuracy',
        n_jobs=-1,
        verbose=1
    )
    grid_search.fit(X_train_scaled, y_train)

    print("Optimal parameters: ", grid_search.best_params_)
    print("Cross-validation accuracy: {:.4f}".format(grid_search.best_score_))

    best_log_reg = grid_search.best_estimator_
    y_pred = best_log_reg.predict(X_test_scaled)
    y_pred2 = best_log_reg.predict(X_train_scaled)

    accuracy = accuracy_score(y_test, y_pred)
    train_accuracy = accuracy_score(y_train, y_pred2)
    cm = confusion_matrix(y_test, y_pred)
    report = classification_report(y_test, y_pred, target_names=['No Heart Disease', 'Heart Disease'])

    print(f"\nTest set accuracy: {accuracy:.4f}")
    print(f"Train set accuracy: {train_accuracy:.4f}")
    print("Confusion Matrix:")
    print(cm)
    print("Classification Report:")
    print(report)

```

Fitting 5 folds for each of 14 candidates, totalling 70 fits

Optimal parameters: {'C': 0.01, 'penalty': 'l2'}

Cross-validation accuracy: 0.8211

Test set accuracy: 0.9000

Train set accuracy: 0.8309

Confusion Matrix:

```
[[47  2]
```

```
[ 7 34]]
```

Classification Report:

	precision	recall	f1-score	support
No Heart Disease	0.87	0.96	0.91	49
Heart Disease	0.94	0.83	0.88	41
accuracy			0.90	90
macro avg	0.91	0.89	0.90	90
weighted avg	0.90	0.90	0.90	90

Compared to our results in (13), the new logistic regression model achieved the highest overall accuracy (0.90) and F1 score (0.88) among all classifiers. It also showed improvement in precision and maintained a relatively similar recall for heart disease prediction, even though those were not part of our initial metrics. This suggests that the added cluster feature was very helpful in capturing subgroup patterns. Overall, this model has the most well-rounded performance and it is a very smart approach.

16)

Team Contributions:

- JC: Questions 1-6
- Jasmine: Questions 7-11 and 15
- Zhiyan: Questions 12-14 and 15

17)

https://github.com/j-asmineho/STATS_3DA3_A6

References

- Cross Validated. (2012, October 30). Clustering as a means of splitting up data for logistic regression [Online forum post]. Stack Exchange. <https://stats.stackexchange.com/questions/31440/clustering-as-a-means-of-splitting-up-data-for-logistic-regression>
- Saxena, A., Minton, D., Lee, D.-C., Sui, X., Fayad, R., Lavie, C. J., & Blair, S. N. (2013). Protective role of resting heart rate on all-cause and cardiovascular disease mortality. *Mayo Clinic Proceedings*, 88(12), 1420–1426. <https://doi.org/10.1016/j.mayocp.2013.09.011>
- Stack Overflow. (2014, February 16). Fine-tuning parameters in logistic regression. Stack Overflow. <https://stackoverflow.com/questions/21816346/fine-tuning-parameters-in-logistic-regression>
- Stack Overflow. (2016, July 17). Random Forest Hyperparameter Tuning - scikit-learn using GridSearchCV. Stack Overflow. <https://stackoverflow.com/questions/35164310/random-forest-hyperparameter-tuning-scikit-learn-using-gridsearchcv>
- Suman, S., Pravalika, J., Manjula, P., & Farooq, U. (2023). Gender and CVD—Does it really matter? *Current Problems in Cardiology*, 48(5), 101604. <https://doi.org/10.1016/j.cpcardiol.2023.101604>